Independent HGTNet: calibration and fresh refit

(a) Calibration loss: 5,040 train / 1,260 validation

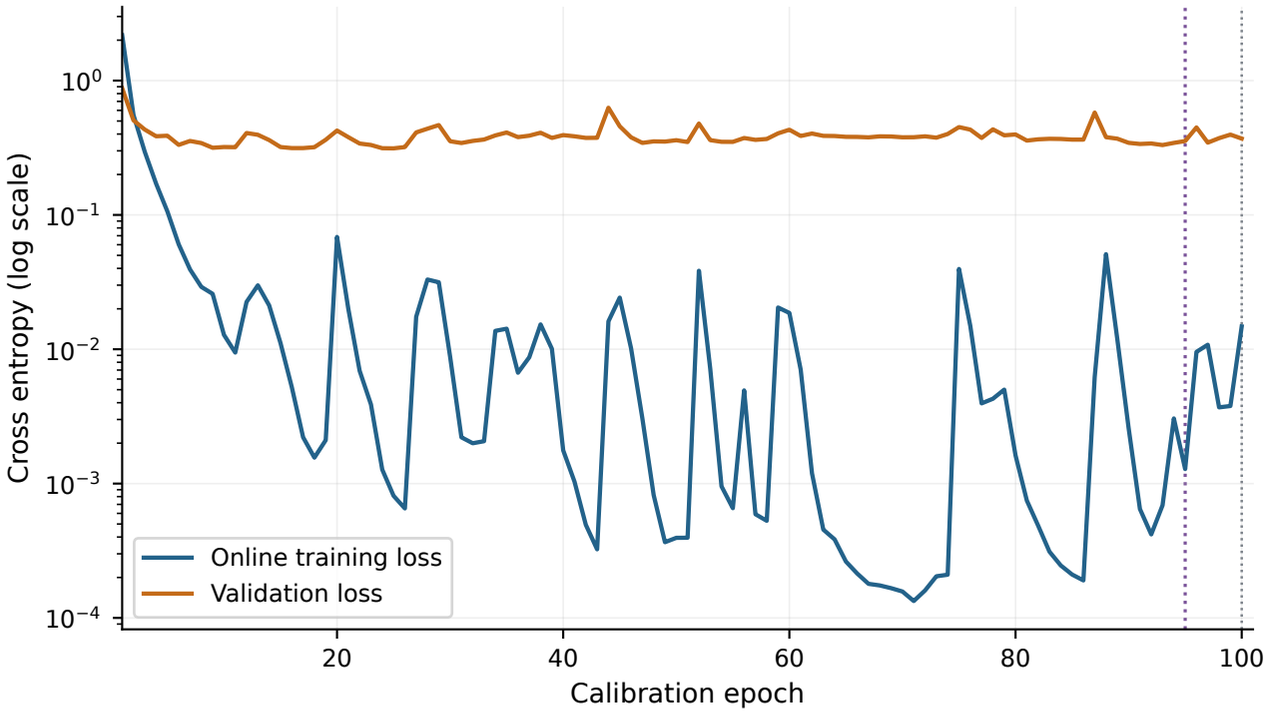

(b) Calibration: validation selects the epoch budget

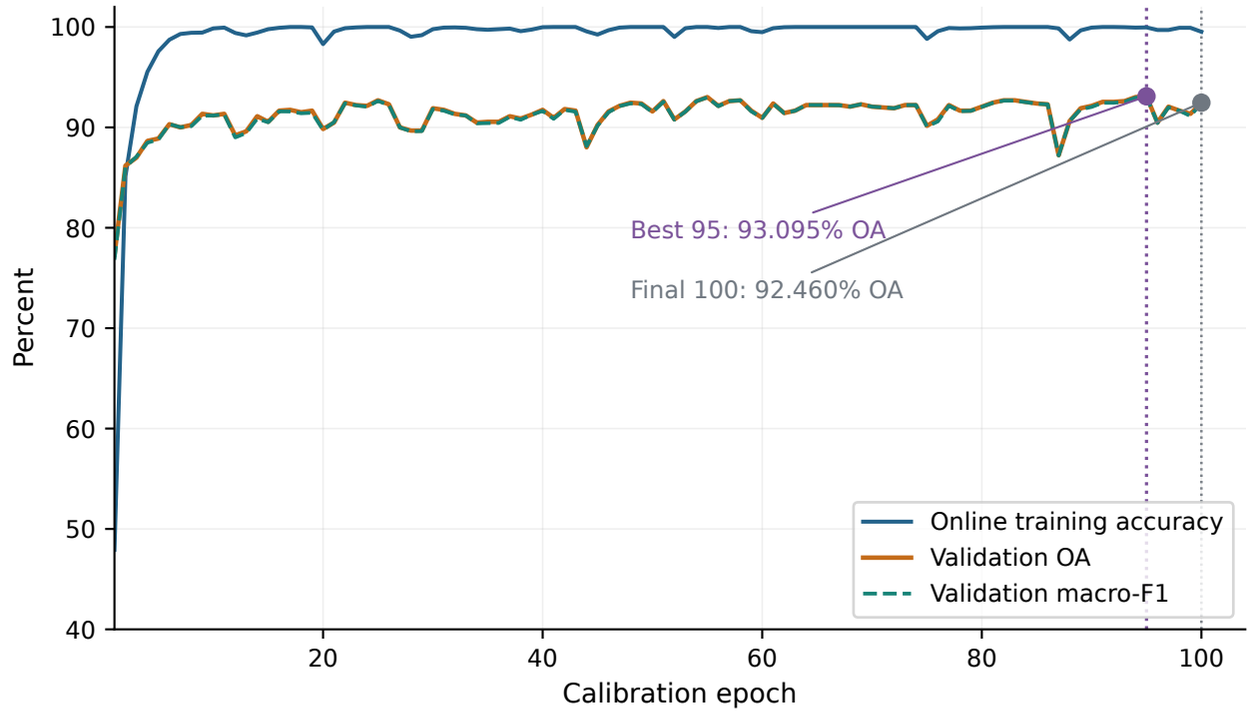

(c) Fresh refit loss: 6,300 pooled train images

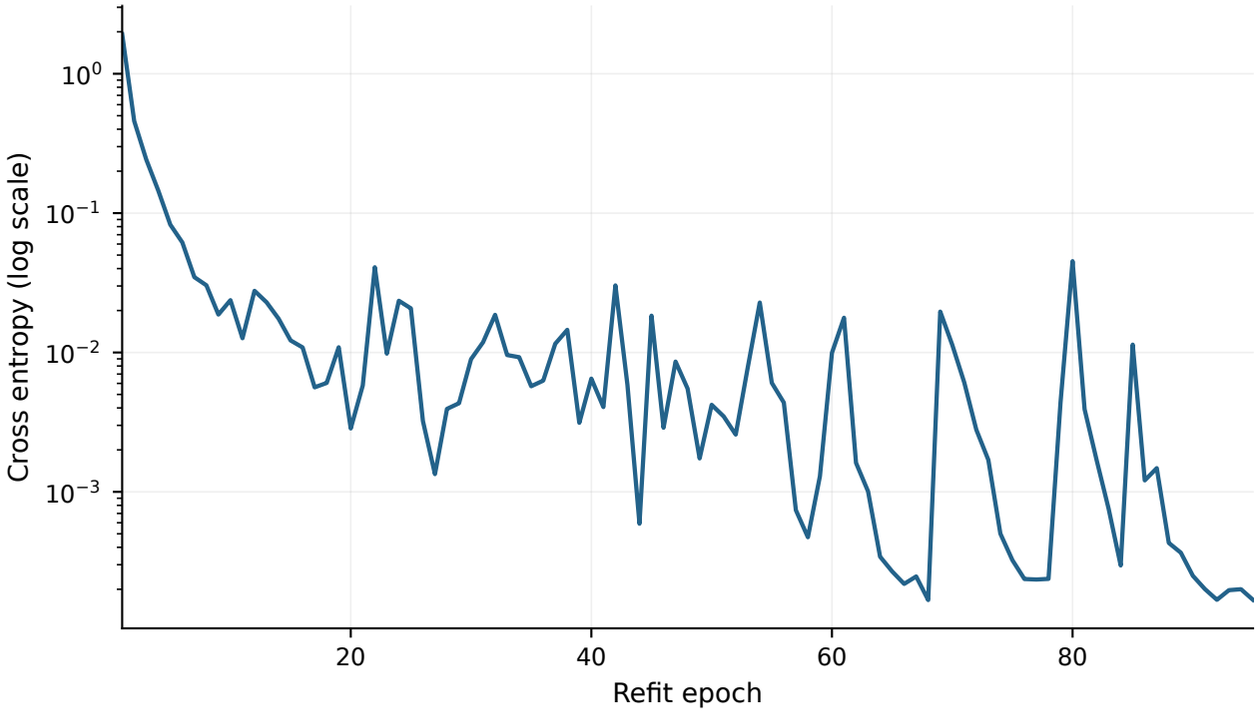

(d) Fresh refit: fixed 95-epoch budget

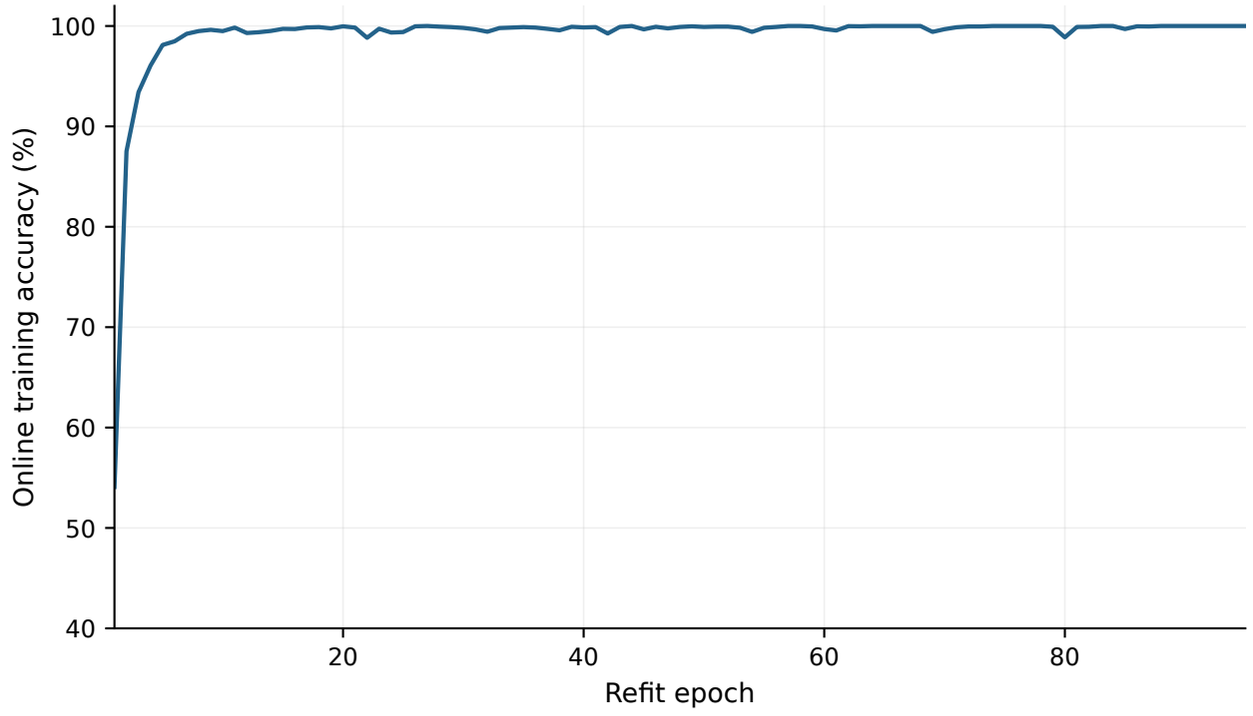

Training accuracy is accumulated while weights update; validation uses the epoch-end model.
Refit starts from the same initial weights and has no validation curve; final test OA = 94.119%, macro-F1 = 94.105%.